

Bachelor's Thesis

COMMIT-FEATURE INTERACTIONS: ANALYZING STRUCTURAL AND DATAFLOW RELATIONS BETWEEN COMMITTS AND FEATURES

SIMON STEUER

November 20, 2023

Advisor:

Sebastian Böhm Chair of Software Engineering

Examiners:

Prof. Dr. Sven Apel

Chair of Software Engineering

Andreas Zeller

CISPA Helmholtz Center for Information Security

Chair of Software Engineering
Saarland Informatics Campus
Saarland University



UNIVERSITÄT
DES
SAARLANDES

Erklärung

Ich erkläre hiermit, dass ich die vorliegende Arbeit selbständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel verwendet habe.

Statement

I hereby confirm that I have written this thesis on my own and that I have not used any other media or materials than the ones referred to in this thesis

Einverständniserklärung

Ich bin damit einverstanden, dass meine (bestandene) Arbeit in beiden Versionen in die Bibliothek der Informatik aufgenommen und damit veröffentlicht wird.

Declaration of Consent

I agree to make both versions of my thesis (with a passing grade) accessible to the public by having them added to the library of the Computer Science Department.

Saarbrücken, _____
(Datum/Date)

(Unterschrift/Signature)

ABSTRACT

Short summary of the contents in English...a great guide by Kent Beck how to write good abstracts can be found here:

<https://plg.uwaterloo.ca/~migod/research/beck00PSLA.html>

CONTENTS

1	Introduction	1
1.1	Goal of this Thesis	1
1.2	Overview	1
2	Background	3
2.1	Code Regions	3
2.2	Interaction Analysis	4
3	Commit-Feature Interactions	7
3.1	Structural CFIs	7
3.2	Dataflow-based CFIs	8
3.3	Combination of CFIs	8
3.4	Feature Size	9
3.5	Feature-Related Concerns of a Commit	9
3.6	Feature Nesting Degree	10
3.7	Implementation	10
4	Methodology	13
4.1	Research Questions	13
4.2	Operationalization	16
5	Evaluation	23
5.1	Results	23
5.2	Discussion	36
5.3	Threats to Validity	39
6	Related Work	41
7	Concluding Remarks	43
7.1	Conclusion	43
7.2	Future Work	43
A	Appendix	45
	Bibliography	47

LIST OF FIGURES

Figure 4.1	Illustration of different commit kinds	16
Figure 5.1	Feature Structural CFIs Plot	25
Figure 5.2	Commit Structural CFIs Plot	26
Figure 5.3	Proportional Dataflow-Plot for Commits	30
Figure 5.4	Proportional Dataflow-Plot for Features	31
Figure 5.5	Author CFI Plot	35

LIST OF TABLES

Table 5.1	Additional Information for Structural Analysis of Features	24
Table 5.2	Additional Information for Structural Analysis of Commits	27
Table 5.3	Additional Information for Dataflow Analysis of Commits	28
Table 5.4	Commits exclusively/partially interacting with features only through dataflow	29
Table 5.5	Relating Inside Dataflow to Outside Dataflow	30

LISTINGS

Listing 3.1	Illustration of structural and dataflow-based CFIs	8
Listing 3.2	Example use-case for the feature-related concerns of a commit	10

ACRONYMS

INTRODUCTION

1.1 GOAL OF THIS THESIS

1.2 OVERVIEW

BACKGROUND

In this section, we summarize previous research on the topic of code regions and interaction analysis. Thereby we give definitions of terms and discuss concepts that are fundamental to our work. In the next chapter, we use the introduced definitions and concepts to explain and investigate structural as well as dataflow-based commit-feature interactions.

2.1 CODE REGIONS

Software Programs consist of source code lines that are translated into an intermediate representation (IR) upon compilation. IR accurately represents the source-code information and is utilized in many code improvement and transformation techniques such as code-optimization. Furthermore, static program analyses are conducted on top of IR or some sort of representation using IR.

Code regions are comprised of IR instructions that are consecutive in their control-flow. They are used to represent abstract entities of a software project, such as commits and features. For this, each code region carries some kind of variability information, detailing its meaning. It should be noted that there can be several code regions with the same meaning scattered across the program. We discuss two kinds of variability information, namely commit and feature variability information, allowing us to define two separate code regions.

Commits are used within a version control system to represent the latest source code changes in its respective repository. Inside a repository revision, a commit encompasses all source code lines that were added or last changed by it.

Definition 1 (commit regions). The set of all consecutive instructions, that stem from source code lines belonging to a commit, is called a *commit region*.

As the source code lines changed by a commit are not necessarily contiguous, a commit can have many commit regions inside a program. To properly address the entire set of these commit regions within a program, we introduce the term regions of a commit. We say that the *regions of a commit* are the set of all commit regions that stem from source code lines of the commit.

In general, features are program parts implementing specific functionality. In this work we focus on features that are modelled with the help of configuration variables that specify whether the functionality of a feature should be active or not. Inside a program, configuration variables decide whether instructions, performing a feature's intended functionality, get executed. Detecting these control-flow dependencies is achieved by deploying extended static taint analyses, such as Lotrack[3].

Definition 2 (feature regions). The set of all consecutive instructions, whose execution depends on a configuration variable belonging to a feature, is called a *feature region*.

Similarly to commits, features can have many feature regions inside a program as the instructions implementing their functionalities can be scattered across the program. Here, we also say that the *regions of a feature* are the set of all feature regions that stem from a feature’s configuration variable.

2.2 INTERACTION ANALYSIS

In the previous chapter, we have already shown how different abstract entities of a software project, such as features and commits, can be assigned concrete representations within a program. We can use interaction analyses to infer interactions between them by computing interactions between their concrete representations. This can improve our understanding of them and give us insights on how they are used inside software projects. An important component of computing interactions between code regions is the concept of interaction relations between them introduced by Sattler [4]. As we investigate structural and dataflow-based interactions, we make use of structural interaction (\odot) and dataflow interaction (\rightsquigarrow) relations in this work.

The structural interaction relation between two code regions r_1 and r_2 is defined as follows:

Definition 3 (structural interaction relation). The interaction relation $r_1 \odot r_2$ evaluates to true if at least one instruction that is part of r_1 is also part of r_2 .

The dataflow interaction relation between two code regions r_1 and r_2 is defined as:

Definition 4 (dataflow interaction relation). The interaction relation $r_1 \rightsquigarrow r_2$ evaluates to true if the data produced by at least one instruction i that is part of r_1 flows as input to an instruction i' that is part of r_2 .

Essential to the research conducted in this paper is the interaction analysis created by Sattler [4]. Their interaction analysis tool VaRA is implemented on top of LLVM and PhASAR. In their novel approach SEAL[6], VaRA is used to determine dataflow interactions between commits. This is accomplished by computing dataflow interactions between the respective regions of commits. Their approach for this will be shortly discussed here, however we advice their paper for a more thorough explanation.

The first step of their approach is to annotate code by mapping information about its regions to the compiler’s intermediate representation. This information is added to the LLVM-IR instructions during the construction of the IR. In their paper, Sattler et al. [6] focus on commit regions, which contain information about the commit’s hash and its respective repository. The commit region of an instruction is extracted from the commit that last changed the source code line the instruction stems from. Determining said commit is accomplished by accessing repository meta-data.

The second analysis step involves the actual computation of the interactions. For this, Sattler et al. implemented a special, inter-procedural taint analysis. It’s able to track data flows between the commit regions of a given target program. For this, information about which commit region affected it, is mapped onto data and tracked along program flow. In this context, we would like to introduce the concept of taints, specifically commit taints. Taints are used to give information about which code regions have affected an instruction through dataflow. For example, an instruction is tainted by a commit if its taint stems from a commit region. It

follows, that two commits, with their respective commit regions r_1 and r_2 , interact through dataflow, when an instruction is part of one's commit region, r_2 , while being tainted by r_1 . Consequently data produced by the commit region r_1 flows as input to an instruction within the commit region r_2 , matching the [dataflow interaction relation](#).

COMMIT-FEATURE INTERACTIONS

In this section, we define structural and dataflow-based commit-feature interactions as well as properties related to them. Furthermore, their meaning and relationship inside a software project is explained here. In the [Overview](#) chapter, we discussed what purpose commits and features serve in a software project. Commits are used to add new changes, whereas features are cohesive entities in a program implementing a specific functionality. In this work, structural interactions are used to investigate how commits implement features and their functionality. In addition, dataflow interactions are examined to gain additional knowledge on how new changes to a program, in the form of commits, affect features.

3.1 STRUCTURAL CFIS

In the background chapter, we discussed the concept of [Code Regions](#), especially commit and feature regions. Logically, we speak of structural interactions between features and commits when their respective regions structurally interact. This structural interaction between code regions occurs when at least one instruction is part of both regions. This is the case when code regions structurally interact through the [structural interaction relation](#) (\odot).

Definition 5. A commit C with its commit regions r_{1C}, r_{2C}, \dots and a feature F with its feature regions r_{1F}, r_{2F}, \dots structurally interact, if at least one commit region r_{iC} and one feature region r_{jF} structurally interact with each other, i.e. $r_{iC} \odot r_{jF}$.

Structural CFIs carry an important meaning, namely that the commit of the interaction was used to implement or change functionality of the feature of said interaction. This can be seen when looking at an instruction accounting for a structural interaction, as it is both part of a commit as well as a feature region. From the definition of [commit regions](#), it follows that the instruction stems from a source-code line that was last changed by the region's respective commit. From the definition of [feature regions](#), we also know that the instruction implements functionality of the feature region's respective feature. Thus, the commit of a structural interaction was used to extend or change the code implementing the feature of that interaction. Following this, we can say that the commits, a feature structurally interacts with, implement the entire functionality of the feature. That is because each source-code line of a git-repository was introduced by a commit and can only belong to a single commit. Thus every instruction, including those part of feature regions, is annotated by exactly one commit region.

Knowing the commits used to implement a feature allows us to determine the authors that developed it. This is made possible by simply linking the commits, a feature structurally interacts with, to their respective authors. By determining the authors of a feature, we can achieve a deeper insight into its development than solely focusing on the commits that implemented it.

3.2 DATAFLOW-BASED CFIS

Determining which commits affect a feature through dataflow can reveal additional interactions between commits and features that cannot be discovered with a structural analysis. Especially dataflows that span over multiple files and many lines of code might be difficult for a programmer to be aware of. Employing VaRA's dataflow analysis, that is discussed in section 3.7, facilitates the detection of these dataflow interactions.

Commit interactions based on dataflow were explained in the [Interaction Analysis](#) section and can be considered as precursors to dataflow-based commit-feature interactions. Similarly to commits interacting with other commits through dataflow, commits interact with features through dataflow, when there exists dataflow from a commit to a feature region. This means that data allocated or changed within a commit region flows as input to an instruction located inside a feature region. This pattern can also be matched to the [dataflow interaction relation](#) (\rightsquigarrow) when defining dataflow-based commit-feature interactions.

Definition 6. A commit C with its commit regions r_{1C}, r_{2C}, \dots and a feature F with its feature regions r_{1F}, r_{2F}, \dots interact through dataflow, if at least one commit region r_{iC} interacts with a feature region r_{jF} through dataflow, i.e. $r_{iC} \rightsquigarrow r_{jF}$.

1. <code>int calc(int val) {</code>	▷ d93df4a	
2. <code>int ret = val + 5;</code>	▷ 7edb283	
3. <code>if (FeatureDouble) {</code>	▷ fc3a17d	▷ FeatureDouble
4. <code>ret = ret * 2;</code>	▷ fc3a17d	▷ FeatureDouble
5. <code>}</code>	▷ fc3a17d	▷ FeatureDouble
6. <code>return ret;</code>	▷ d93df4a	
7. <code>}</code>	▷ d93df4a	

Listing 3.1: Illustration of structural and dataflow-based CFIs

The above code snippet contains both structural as well as dataflow-based commit-feature interactions. Commit `fc3a17d` implements the functionality of `FeatureDouble` for this function. It follows that a structural commit-feature interaction can be found between them, as their respective commit and feature regions structurally interact. Commit `7edb283` introduces the variable `ret` that is later used inside the feature region of `FeatureDouble`. This accounts for a commit-feature interaction through dataflow, as data that was produced within a commit region is used as input by an instruction belonging to a feature region of `FeatureDouble` later on in the program.

3.3 COMBINATION OF CFIS

When investigating dataflow-based CFIs, it is important to be aware of the fact that structural CFIs heavily coincide with them. This means that whenever commits and features structurally interact, they are likely to interact through dataflow as well. As our structural analysis has already discovered that these commits and features interact with each other, we are more interested in commit-feature interactions that can only be detected by a dataflow analysis. That this relationship between structural and dataflow interactions exists becomes clear when looking at an instruction accounting for a structural CFI. From definition 5, we know that the instruction belongs to a commit region of the interaction's respective commit. It follows that

data changed inside the instruction produces commit taints for instructions that use the data as input. Now, if instructions that use the data as input are also part of a feature region of the interaction's respective feature, the commit and feature of the structural interaction will also interact through dataflow. However, such dataflow is very likely to occur, as features are functional units, whose instructions build and depend upon each other. Knowing this we can differentiate between dataflow-based CFIs that occur within the regions of a feature and those where data flows from outside the regions of a feature into them. From prior explanations, it follows that this differentiation can be accomplished by simply checking whether a commit that influences a feature through dataflow, also structurally interacts with it.

3.4 FEATURE SIZE

When examining commit-feature interactions in a project, it is helpful to have a measure that can estimate the size of a feature. We can use such a measure to compare features with each other and, thus, put the number of their interactions into perspective. Considering our implementation, it makes most sense to define the size of a feature as the number of instructions implementing its functionality inside a program. As the instructions inside the regions of a feature implement its functionality, we can define the size of a feature as follows:

Definition 7. The *size* of a feature is the number of instructions that are part of its feature regions.

3.5 FEATURE-RELATED CONCERNS OF A COMMIT

Often, multiple features are implemented in the same code space. This can occur due to various reasons, for example being a sub-feature of an overarching feature or depending on data generated by other features. Logically, the instructions stemming from the respective code are part of several feature regions. At the same time, only a subset of these instructions actually implement functionality of a specific feature. We say that these features are, at least partially, nested features. If a developer introduces changes to such a nested feature, the according commit will also produce structural CFIs with other features. Now, we are faced with the issue that, while the commit structurally interacts with several features, it only changes functionality of one. In this case, it is impossible to tell which features the commit really changed from the raw data of its structural CFIs alone. Therefore, we say that the commit only deals with a single feature-related concern as long as the code its interactions stem from belong to the same set of features. This code-property indicates that a single nested feature was changed by the commit instead of multiple features at once. In our analysis, we are working with instructions and therefore speak of the same code space for features when the instructions we are dealing with belong to regions of the same features. It follows that we can assign a commit a number of concerns according to its structural CFIs and the instructions they occur in:

Definition 8. The *feature-related concerns* of a commit are the different sets of features the commit structurally interacts with at the same time, i.e. in the same instructions.

They should be used as a measure to gauge the responsibilities of a commit that are difficult to decipher due to feature-nesting. An exemplary use-case is shown in the listing below.

26. <code>if</code> (Min) {	▷ b9ea2f3	▷ Min
27. <code>val</code> = min(x,y)	▷ b9ea2f3	▷ Min
28. <code>if</code> (CheckZero) {	▷ e37a1c8	▷ Min ∧ CheckZero
29. <code>if</code> (val == 0) {	▷ e37a1c8	▷ Min ∧ CheckZero
30. <code>throw</code> Error();	▷ e37a1c8	▷ Min ∧ CheckZero
31. }	▷ e37a1c8	▷ Min ∧ CheckZero
32. }	▷ e37a1c8	▷ Min ∧ CheckZero
33. }	▷ b9ea2f3	▷ Min

Listing 3.2: Example use-case for the feature-related concerns of a commit

The above code-snippet illustrates how considering the feature-related concerns of a commit is more insightful than simply looking at the features it structurally interacts with.

We can see that the commit e37a1c8 exclusively implements functionality of the CheckZero-feature. Still, it also structurally interacts with the Min-feature, as the code of CheckZero is nested inside of it. However, commit e37a1c8 only has a single feature-related concern, as it only affects source-code lines that encompass regions of the same set of features. If commit b9ea2f3 instead introduced the entire code-snippet, it would have two feature-related concerns. In both cases, the feature-related concerns accurately reflect the number of features that were implemented by a commit. At the same time, this is not true for the number of features a commit structurally interacts with.

3.6 FEATURE NESTING DEGREE

As dicussed in section 3.5, feature-nesting is an important issue worth taking into account. Instructions that are part of several feature regions will, by definition, also consitute the size of the respective features. At the same time, it is unlikely that the instruction actually implements functionality for every one of them. Thus, the size of a feature can encompass many instructions not implementing its functionality. Technically, we can only be certain of an instruction's purpose if it is exclusively part of a single feature region. In order to achieve a more accurate description of a feature's size, we introduce the concept of a feature's nesting degree:

Definition 9. The *nesting degree* of a feature is the fraction of its size where its respective regions do not appear exclusively.

Besides that, the nesting degree of a feature informs us about the frequency of feature-nesting within a feature and, in combination with other features, the entire project.

3.7 IMPLEMENTATION

The detection of structural as well as dataflow-based commit-feature interactions is implemented in VaRA [5]. Additionally to commit regions, VaRA maps information about its feature regions onto the compiler's IR during its construction. Commit regions contain the hash

and repository of their respective commits, whereas feature regions contain the name of the feature they originated from. VaRA also gives us access to every llvm-IR instruction of a program and its attached information. To accommodate later evaluations using nesting degrees, we create reports featuring a more complex datatype than a sole CFI. For this, we specifically save which features a commit structurally interacts with at the same time, i.e. inside the same instructions. We know that an instruction is always part of exactly one commit region, but could possibly belong to any number of feature regions. According to definition 5, we encounter a structural CFI, if an instruction belongs to at least one feature region. For every such instruction, we store the discussed interaction, as the commit and the features present in the instruction. For each interaction, we also save the number of instructions it occurs in. This is accomplished by incrementing its instruction counter if we happen to encounter a duplicate. By creating complex structural reports, we can not only determine the nesting degree of each structural CFI later on, but also the nesting degrees of the instructions they occur in. This allows us to calculate the definite size of a feature, additionally to its normal size defined in definition 7. We can compute the size of a feature as the sum over the instruction counters of all found interactions the feature is part of. To calculate a feature's definite size, we only consider interactions where the feature appears exclusively. These interactions stem from instructions with a nesting degree of one, which means that they are part of no further feature regions.

In the [Interaction Analysis](#) section, we discussed the taint analysis deployed by VaRA. There, VaRA computes information about which code regions have affected an instruction through dataflow. Checking whether a taint stems from a commit region allows us to extract information about which commits have tainted an instruction. Thus, dataflow-based commit-feature interactions can also be collected on instruction level. According to definition 6, we can store a dataflow-based interaction between a commit and a feature, if an instruction has a respective commit taint while belonging to a respective feature region. Consequently said instruction uses data, that was changed by a commit region earlier in the program, as its input.

METHODOLOGY

The purpose of this chapter is to first formulate the research questions that we examine in our work and then propose our method of answering them.

4.1 RESEARCH QUESTIONS

In the [Interaction Analysis](#) chapter we discussed the different meanings of structural and dataflow-based interactions. With this knowledge, we can answer many interesting research topics. These topics include patterns in feature development and usage of commits therein as well as findings about how likely seemingly unrelated commits are to affect features inside a program.

RQ1: How do commits and features structurally interact with each other?

We intend to research two main properties which already provide a lot of insight into the development process of features and best practices of commits therein.

Investigating Patterns around Feature Development

Firstly, we examine the number of commits features interact with structurally. This gives us an estimate on how many commits were used in the development of a feature and the according distribution among the features of a project. Our analysis also allows us to measure the size of feature, which can put the number of commits used to implement a feature into perspective. We refine our investigation by factoring in the nesting degrees of features (9). In sections 3.5 and 3.6, we discussed feature-nesting and the risks it bears for our analysis. Namely that a fraction of a feature's size can consist of instructions that do not implement its functionality. The same is the case for commits structurally interacting with a feature, meaning that a portion of them possibly didn't participate in its development, but in the development of other features. We therefore consider the frequency of feature-nesting within a project, so that we can tell how accurately our initial data reflects the development of features.

Examining the Usage of Commits in Feature Development

Secondly, we investigate the number of features a commit usually changes by examining its structural CFIs. This is especially interesting when considering best practices surrounding the usage of commits. It is preferred to keep commits atomic[1] meaning they should only have a single responsibility. As different features implement separate functionalities, it's unlikely for a commit to change several features while dealing with the same issue. Transferring this to our work, high quality commits should mostly change a single feature. Acquiring data on this issue might show how strictly this policy is enforced in the development of features across

different projects. In section 3.5, we discussed that, due to feature nesting, a commit can structurally interact with multiple features at the same time while only changing functionality of one of them. Instead of considering the features a commit structurally interacts, we rather examine the number of its feature-related concerns (8) for this RQ. This way we aim to achieve a more accurate prediction for the number of features a commit changed or implemented. In our analysis, we should ideally filter commits whose purpose was refactoring code, since they do not fit our criteria of changing or adding functionality to a feature. Exceptionally large commits might be best considered for this, as this is where we expect most commits used for large-scale refactoring to appear. A qualitative review of these commits is still necessary to make a final decision on whether they should be filtered.

RQ2: How do commits interact with features through dataflow?

Investigating dataflow inside a program can unveil interactions between its entities that were previously hidden from programmers. This can help them understand the extent to which different parts of a program influence each other. Previous studies have laid the groundwork for researching dataflow interactions between different abstract entities of a program. While it has shown a wide range of interesting use-cases, it has focused solely on dataflow between commits. That's why we aim to provide first insights into the properties of dataflow-based CFIs.

The Proportion and Dependencies of Commits Affecting Features through Dataflow

Firstly, we investigate how connected commits and features are by analyzing the number of features a commit affects through dataflow. Knowing what fraction of all commits contributing code to a project are part of dataflow-based interactions can show how often new commits affect the data of a feature. Regarding this, it is worth considering the dependency discussed in section 3.3, that structural interactions heavily coincide with dataflow interactions. This implies that commits constituting code of a feature are very likely to influence said feature through dataflow as well. In section 3.3, we also mention that the dataflow of commits not structurally interacting with a feature, must stem from outside the regions of the feature. Naturally, said dataflow is less intentional and subsequently more interesting, than dataflow occurring inside the regions of a feature. Programmers are less aware that changes introduced with these commits might affect the data of seemingly unrelated features. Therefore, we intend to differentiate between dataflow interactions with an outside and those with an inside origin in our analysis. This allows us to especially focus on commits part of outside dataflow interactions and examine their proportion among all commits.

Understanding Features and the Commits Affecting them Through Dataflow

Dataflow-based CFIs allow us to examine another interesting property, namely the number of commits affecting a feature through dataflow. Here, we differentiate between commits `INSIDE` and `OUTSIDE` of features, i.e. commits that either do or don't contribute code to them. We examine the ratio of outside to inside commits influencing a feature to gauge where most dataflow interactions of a feature stem from. Considering our previous assumptions, the ratio

of outside to inside commits could decrease with increasing size of a feature. Smaller features are likely to structurally interact with less commits than their larger counterparts, resulting in them also interacting with less inside commits through dataflow. However, the number of outside commits affecting data of a feature might not be dependent on feature size to such an extent. This leads us to another relationship worth investigating, namely the relation between the size of a feature and the number of outside and inside commits interacting with it through dataflow. We examine the two commit kinds separately, as we already have a strong supposition for inside commits, but are less certain for outside commits. Determining to what extent feature size is the driving factor in this relation, could tell us whether it is worth considering other possible properties of features in future analyses. Discovering features with unusually many outside commits in relation to their size could also allow us to detect feature-code that is central-code for a project.

RQ3: How do authors interact with features?

A simple yet promising way to extract more information out of the collected CFIs is to link the interaction's commits to their respective authors. Thus, we are able to investigate how authors interact with features both structurally and through dataflow. Similarly to the number of commits implementing a feature, we can now calculate the number of authors that participated in its development. The same applies to commits interacting with features through dataflow, for which we can now determine the authors contributing these commits. Here, we only consider outside commits, i.e. commits not constituting code of the feature, since we have already considered all inside commits when investigating the developers implementing a feature. We are especially interested in how many authors exclusively interact with a feature through dataflow in comparison to the number of its developers. This lets us determine whether the developers that implement features are also the ones mainly responsible for changes affecting them through dataflow.

Finally, we relate both types of author interactions to the respective size of a feature. Regarding structural interactions, this allows us to determine whether a feature's extent in the source-code of a project is driving factor in the amount of authors needed to implement it. Software companies could use evidence on this issue as advice on how to allocate programmers on to-be implemented features. Findings on the investigated dataflow interactions of authors could tell developers that their changes might affect small and, at first sight negligible, features surprisingly often. We also have a look at the specific functionality of features since their purpose could also play an important role in the number of authors interacting with them.




Figure 4.1: Illustration of different commit kinds

In the first two **RQs** we have discussed different kinds of commits and the ways in which they interact with features. The above figure showcases them in a venn diagram and illustrates the dependencies and divisions between them.

4.2 OPERATIONALIZATION

Here, we explain how the proposed RQs can be answered. The general experiment process is the same for all RQs. At first, we collect data comprising all structural or dataflow-based CFIs by creating reports of a specific type for a chosen software project. The collected data is then processed in order to gain information for each commit or feature in the project, such as the number of interacting commits and size of a feature. The processed data is used to calculate statistical information, such as the mean and variance, or the strength of a correlation. To facilitate a faster and better understanding of the processed data and calculated statistics, we display them graphically via bar or regression plots. The projects we investigate in this work, for example `xz` and `gzip`, are of a small size and are used in a compression domain. This choice is based on the fact, that our research intends to only lay basic groundwork, where smaller projects can already offer a lot of insight. Since we only investigate a few projects, we chose them to be of a similar domain, such that a comparison between them makes more sense. In the following sections, we explain our method of investigation in more detail for each RQ.

RQ1: How do commits and features structurally interact with each other?

For this RQ, we examine the interactions contained in the structural reports created with our [Implementation](#) in VaRA. As there could be vast differences between the projects, we also work with and present their respective data separately.

Investigating Patterns around Feature Development

From the collected structural reports, we first extract the number of structurally interacting commits and size for each feature. The datapoints of the two properties are shown in two

respective bar plots, where the values for each feature are shown in a separate bar. The bars are labelled after the name of the feature they present, to allow for comparisons between the two plots and between different projects. For comparisons between the investigated projects, we also calculate the average number of interacting commits and size of a feature for each project. For a simpler recognition of the existing distribution, the features are sorted in an increasing order according to their value in the metric of the respective plot. This also allows us to see whether the rank of a feature in one metric is indicative of its rank in the other metric and vice versa. Additionally to the size of a feature, we also determine its nesting degree(9). The average nesting degree among the features of a project is subsequently shown in a \LaTeX -table. The strength of a project's overall nesting degree can in turn give us a direct estimate on the frequency of feature-nesting inside of it. Finally, we relate the two investigated properties in a `SEABORN` regression plot. Here, we show whether features increasing in size, in turn structurally interact with more commits, i.e. need more commits to be implemented. Therefore, the value of the x-axis shows the size of a feature, whereas the y-axis shows the number of interacting commits. Each feature is represented by a scatter point inside the plot combining its respective y-values from the two previous plots. A linear regression line is drawn in the plot matching the occurring scatter points, where a rising graph could already indicate a positive correlation. To check whether we are dealing with a statistically significant linear correlation, we compute the pearson correlation coefficient and its p-value in the `STATS` submodule of the `SCIPY` python namespace. Both values are subsequently inserted into the regression plot with the text-function of the `MATPLOTLIB`. To confirm our initial suspicion of a strong positive correlation, the according correlation coefficient must be close to one, while the p-value must be close to zero falling below our rejection interval of 95%.

Examining the Usage of Commits in Feature Development

For each commit part of at least one structural interaction, we determine the number of its feature-related concerns (8). For this, we examine the complex structural CFIs discussed in our [Implementation](#). There, we increment the number of a commit's concerns upon finding a new set of features it structurally interacts with at the same time, i.e. within the same instructions. Following this, we give a comprehensive overview of our results in a `SEABORN` histplot for every investigated project. The x-axis of the histogram shows the number of a commit's concerns, whereas the y-axis shows the number of commits for which the respective x-value matches. If 50 commits deal with a single concern inside a project, the y-value of one will be 50 accordingly. Thus, we can quickly see to what extent our derived best practice surrounding commits in feature development is enforced in a project. The more commits are located at one in comparison to other x-values, the more commits are likely to have only changed a single feature. To facilitate comparisons between projects, we choose the same x- and y-labels for every histogram. Commits with exceptionally many feature-related concerns are shown in an additional x-tick to avoid an overly long x-axis. Instead, we mention these commits explicitly during our evaluation and relate the number of their concerns to their specific purpose.

In a second analysis step, we filter exceptionally large commits, whose purpose we expect to be something else than implementing features, such as code-refactoring. Similarly to the regions of a commit, we determine its size in a repository as the number of source-code lines

that were last changed or added by it. We then apply tukey’s fence with a k-value of 3 to filter far outliers of commit sizes among commits interacting with features within their respective project. Afterwards, we perform a qualitative review of the filtered commits to check whether our initial suspicion about their purpose is correct. To assess the effect of filtering exceptionally large commits, we compute the average number of concerns within a project before and after the application of our restriction. Both averages are subsequently shown in a cross-project table so that we can compare them within a project and recognize a general trend among them.

RQ2: How do commits interact with features through dataflow?

The projects investigated for dataflow-based CFIs are the same projects as investigated for structural CFIs. This choice gives us more insight into a single project and allows us to combine both analysis results as will be discussed below.

Analyzing the Proportion and Dependencies of Commits Affecting Features through Dataflow

We do not consider LRZIP here, as we are only able to detect regions of three out of 10 features in the project. This means that we can only find a fraction of all commits interacting with features, resulting in significantly lower determined proportions.

An initial step of evaluating what fraction of commits affect features through dataflow, is determining which set of commits to consider in the first place. Logically, we should only consider commits that could potentially be part of a dataflow-based CFI. The only prerequisite for this is that commits are represented by commit regions inside a program. This is the case when there exists at least one source-code line that was last changed or added by them. We call the respective commits, active commits as their contributed code is still “active” in the repository. Following the calculation of the number of active commits in the projects, we begin examining the created dataflow reports. For each commit part of at least one dataflow-based CFI, we save the set of features it interacts with. By dividing the number of these commits by the number of all active commits, we can determine their proportion within their respective project. To allow for an overview of all projects, we show the calculated percentages in a seaborn bar plot. To provide evidence for our claim that structural CFIs heavily coincide with dataflow-based CFIs, we now compute the proportion of commits with dataflow interactions among the commits with structural interactions. Given that the latter is significantly higher than the overall proportion of commits with dataflow interactions, our initial notion will be confirmed. Alongside the percentage of commits part of dataflow-based CFIs, we also show the percentage of active commits that are part of structural CFIs in the first bar plot. As we expect most of the latter commits to be part of dataflow-based CFIs as well, the difference between their respective percentages indicates what proportion of commits interact with features exclusively through dataflow. Following the broad overview of commits affecting features through dataflow, we intend to specify and differentiate between the origin of said dataflow. For each feature a commit interacts with through dataflow, we check whether they also structurally interact with each other. If they do structurally interact, we speak of an inside dataflow origin, otherwise we speak of an outside origin. Naturally, dataflow occurring inside the regions of a feature is more intentional and therefore less interesting,

than dataflow originating outside its regions. In section 3.3, we explain, that the dataflow of commits and features not structurally interacting with each other, must be outside dataflow. Even though the regions of a commit and regions of a feature must partially overlap in order for them to structurally interact, such a commit can still have regions not part of any feature regions. Therefore, we cannot be sure whether the dataflow of these interactions originates from outside or inside the regions of a feature. Due to the inherent properties of structural interactions, i.e. heavily coinciding with dataflow interactions, we assume their dataflow to originate inside the regions of a feature. Upon determining the dataflow origin of each CFI, we split the commits with dataflow interactions into three categories. Next to commits that only affect features either through inside or through outside dataflow, a commit can also interact with multiple features, once through inside and once through outside dataflow. For each project, we calculate the respective percentages of each category and present them in a stacked bar plot. Logically, each bar representing a project adds up to 100% as every commit falls into exactly one category. The commits falling into categories of at least partial outside dataflow origins are part of interactions, which we could only detect with our dataflow analysis. To quantify the according commits within a project, we calculate their number and percentage among a project's active commits. Both bar plots are displayed next to each other in the same figure, where the sequence in which the projects are shown in both plots is determined by the proportion of commits with dataflow interactions. This allows for an easier comparison within a project, as its respective bars will have the same position in both plots. Statistical values, including the number of active commits and the probability for a commit to be part of a dataflow-based CFI given that it structurally interacts with features, are not included in the plots. As they can aid the understanding of the plots and add important context to them, we display the respective values in a \LaTeX -table.

Exploring Features and the Commits Affecting them Through Dataflow

Initially, we determine the set of commits each feature interacts with through dataflow. The according information is extracted from the same dataflow reports used in the previous section of this RQ. Depending on whether the commit also structurally interacts with a feature, we split the set of commits for each feature into two categories. Outside commits are fully located outside the regions of a feature, meaning that their data flows into a feature. Inside commits are at least partially located inside a feature making it very likely that their dataflow occurs within the regions of a feature. Naturally, said dataflow affecting a feature is less interesting than dataflow stemming from seemingly unrelated commits. To gauge what commit kind makes up the majority of commits affecting features through dataflow, we present both values for each feature in a SEABORN bar-plot. For every examined project, we display the number of outside and inside commits with separate bars for the different features. To enable an exact comparison, we also calculate the median for both outside and inside commits of a project. We choose the median and not the mean, as we are especially interested in detecting a general trend across the features of a project and intend to avoid outliers skewing the data in one direction. Besides that, features with an exceptionally high number of interacting commits or a high proportion of commits of a certain kind are already shown in the bar plots. The features are explicitly named to allow for cross-project comparisons of certain types. In a second analysis step, we now relate the size of a feature to the number

of commits affecting it through dataflow. Here, we are interested in determining whether and to what extent there exists a positive linear relationship between the two. We employ two linear regression analyses, relating the size of a feature once to its outside and once to its inside commits. Similarly to our investigation in **RQ1**, the results are displayed in a `SEABORN` reg plot for each project. In each regression plot, we also show the two computed pearson correlation coefficients and their respective p-values. We decide for a 97.5% rejection interval again, meaning that the p-values must be lower than 0.025 to provide enough evidence that the two datasets are indeed not un-correlated. Furthermore, the correlation coefficient must range close to 1 in order to warrant a strong positive correlation. Lastly, we investigate our claim that smaller features have a higher proportion of outside to inside commits affecting it through dataflow. For this, we examine the relation between the size of a feature and the ratio of its outside to inside commits. Again, we create linear regression plots for each project and calculate the according pearson correlation coefficients and the respective p-values. We could also perform other regression analyses, for example a logarithmic regression, as they could fit our datapoints better. Following our explanations in section 4.1, we expect the ratio of outside to inside commits to decrease with increasing feature size. This would result in a negative linear correlation coefficient and a falling linear regression line. It remains to be seen whether the respective values are statistically significant across several projects.

RQ3: How do authors interact with features?

Technically, we perform the same analysis regarding features as discussed for the two previous RQs. The only difference is that we replace the commits interacting with a feature with their respective authors. To facilitate expressing this, we simply say that the respective authors interact with features based on a certain interaction type or that they are a feature's structural or dataflow authors. In section 4.1, we mentioned that we only consider commits outside of features when examining the authors interacting with them through dataflow. The computation of these commits has already been explained in the operationalization of **RQ2**, which we reuse for this RQ. The number of authors that interact with a feature either structurally or through dataflow are presented in a `SEABORN` bar plot. Similarly to the previous RQs, the results of each project are shown in separate plots with labeled ticks for each feature. We also calculate the mean number of dataflow and structural authors of a project to give concrete information from what type of interaction most author interactions stem from. In a third bar for each feature, we display the number of authors that exclusively interact with them through dataflow. There, we exclude authors who also structurally interact with the feature, i.e. the authors making up the first bar. This shows us how many authors affecting a feature can only be determined via our dataflow analysis. For the structural interactions of authors, we compute what fraction of them also affect the respective features through dataflow. This can tell us to what extent the developers implementing a feature are likely to introduce code changes whose data the feature uses later on. Additionally considering whether there are more structural than dataflow authors lets us determine whether the developers of a feature are also the ones mainly responsible for changes affecting them through dataflow. Finally, we perform the same linear regression analysis as we have done for feature-sizes and the number of commits that interacting with them. By calculating the strength of their correlation, we are able to determine whether the presumed extent of a feature inside a project is an accurate

predictor for the number of developers needed to implement it. The authors interacting with a feature through outside dataflow are also indirectly involved in its development, since the data stemming from their contributed commits is necessary for the functionality of the feature.

EVALUATION

This chapter evaluates the thesis core claims. - for this, we investigate four projects

5.1 RESULTS

RQ1: Evaluation of Structural CFIs

Examining the structural CFIs of different projects can give us insights into the patterns of feature development and the usage of commits therein. We first generally describe our results with the help of figure 5.1 and figure 5.2 as well as the tables 5.1 and 5.2. We also explain particularly interesting results of individual commits and features by investigating their specific purpose and properties.

Patterns of Feature Development

In the first part of RQ1, we examine the number of commits involved in the development of a feature and relate it to its size. Figure 5.1 illustrates our results in three different plots for each project. Each row displays the results for one project with the name of the respective project being shown on the far left. In the first column, we show the number of structurally interacting commits for each feature in a bar-plot. For all projects, we notice a wide distribution of structurally interacting commits between features. Numerous features, across all projects, interact with less than 5 commits suggesting that they need very little work to be implemented in comparison to other features. Both xz and LRZIP have features structurally interacting with more than 20 commits, while GZIP even has four features that interact with more than 60 commits. With a mean of 29, the features of GZIP have by far the highest number of interacting commits on average. xz, LRZIP and BZIP2 have much lower averages at 9, 16 and 4 respectively. As all projects are of a compression domain, we can find several features with the same general functionality across the examined projects. The verbosity-feature structurally interacts with the most commits in xz and BZIP2, while the same is not the case for the verbose-feature of GZIP. Although GZIP encompasses the highest average number of interacting commits, its recursive-feature only interacts with 2 commits, while the recursive-feature of LRZIP interacts with 24 commits.

In the second column of figure 5.1, we display the calculated feature-sizes for each project. Again, we see a wide range of different feature-sizes within all projects. Most notably, there are large jumps between adjacent features, for example from to_stdout with a size of 450 to no_name with a size of 7300 in GZIP. Similarly to the number of interacting commits, GZIP has by far the highest average feature size at 2500. The average feature sizes for the other projects range from 190 for xz to 390 for BZIP2. In table 5.1, we display the average nesting degree of features for all projects. Feature-nesting is the least common for xz and LRZIP and most common for GZIP with an average of 0.93 indicating that the majority of features are nested to

Table 5.1: Additional Information for Structural Analysis of Features

Projects	Average Feature Nesting Degree
xz	0.34
gzip	0.93
bzip2	0.64
lrzip	0.37

large extents. This is also the case for the four, by far largest, features of GZIP, namely `no_name`, `method`, `force` and `decompress`. Specifically, 90 to 100% of the instructions constituting their size are identical among the mentioned features. It is therefore not surprising that most, particularly 66, of the commits interacting with each feature are identical for all four features. For them, it is highly doubtful whether the size of a feature and the number of interacting commits reflect the actual number of implementing commits and instructions of a specific feature.

Finally, the metrics used in the two previously discussed plots are compared to each other using the regression plot of figure 5.1. That is, we compare the size of a feature with the number of commits that structurally interact with it. The values of the blue dots each represent a feature and are therefore identical to the values of the bars in the two previous columns. For both xz and GZIP, we determine a strong positive correlation between the size of a feature and its number of structurally interacting commits. Their linear correlation coefficients are close to 1 at 0.91 and 0.978 respectively with p-values smaller than 10^{-4} . This data provides strong evidence that the observed correlation is statistically significant. While BZIP2 and LRZIP have positive correlation coefficients of 0.585 and 0.968 respectively, their p-values are relatively high at over 0.1. The high p-values can be explained by a lack of datapoints for LRZIP and conflicting datapoints for BZIP2. For example, do both the `opMode`- and `srcMode`-feature of BZIP2 structurally interact with 6 commits, but encompass vastly different feature sizes. Overall, the two projects do not produce conclusive statistical evidence that the size of a feature and the number of interacting commits is positively correlated.




Figure 5.1: Feature Structural CFIs Plot




Figure 5.2: Commit Structural CFIs Plot

Usage of Commits in Feature Development

As discussed in chapter 4.1, a best practice for commits in feature development is that one commit should generally change only one feature. We investigate to what extent this holds in our investigated software projects by examining the number of feature-related concerns of a commit. The distribution of how many commits have a certain number of concerns is displayed in figure 5.2. xz, GZIP and LRZIP show similar distributions, which is why we discuss them now and bzip2 later. For xz, GZIP and LRZIP, we note that the majority of commits only have single feature-related concern. The according percentages for commits with a single concern among all considered commits are 72%, 69% and 81%. The commit count gradually drops with an increasing number of concerns until it subsides to 0 at an x-value of 7, 6 and 4 respectively. Interestingly, three commits of GZIP have more than 10 concerns with the commit 33ae4134cc even encompassing 47 concerns. Said commit originally introduced the `main.c`-file containing the `main`-function among other things. The exceptionally high number

Table 5.2: Additional Information for Structural Analysis of Commits

Projects	Average Number of Concerns	After Filtering Large Commits
xz	1.48	1.29
gzip	2.40	1.55
bzip2	5.11	4.62
lrzip	1.30	1.30

of 33ae4134cc's concerns can be explained by the fact that many of GZIP's 14 features are partially implemented within gzip.c. Additionally, the code of features is heavily nested inside each other, which means that the instructions stemming from said code encompass many different combinations of feature regions. For BZIP2, there are much less commits used in feature-development compared to the other projects with only 9 commits part of structural CFIs. The distribution consists of two clusters, the first being at a single concern and the second cluster being located between 7 and 10 concerns. The second cluster encompassing 5 commits is slightly bigger than the first cluster with a commit count of 4. Since we are only dealing with a few commits here, that still produce quite interesting and unique data, we have a closer look at the purpose of these commits and relate it to the number of their concerns. We find that all commits part of the second cluster were used to publish new versions of BZIP2, always encompassing over 1000 additions. Three out of the four commits part of the first cluster introduced minor changes and bug fixes, never changing more than 30 source-code lines. The remaining commit updated BZIP2 to version 1.0.4 fixing minor bugs from the previous version with many of the added source-code lines being used for comments.

In table 5.2, we display the average number of feature-related concerns within a project both before and after filtering exceptionally large commits. The first column contains the project to which the averages in their respective row belong. In the following column, we show the initial averages before applying the mentioned restriction. XZ and LRZIP encompass the lowest average number of concerns at 1.48 and 1.3 respectively. Although the distribution of GZIP shown in figure 5.2 is similar to those of the formerly mentioned projects, its initial average is comparatively high at 2.4. The main cause for this are the discussed three commits of GZIP with exceptionally many concerns. After filtering large commits within a project, we note that the averages of XZ, LRZIP and GZIP become more similar. While the average slightly decreased for XZ to 1.29, it stayed the same for LRZIP. The three discussed commits of GZIP were part of the filtered commits, which explains the large drop of its average to 1.55. Unsurprisingly, the initial average of BZIP2 is by far the highest among all projects. Since most of BZIP2's commits are relatively large, this is also the reason why they are not excluded from our analysis. Here, only one commit with 9 concerns is removed, which lowers the average to 4.58. We also performed a qualitative analysis of the filtered commits to check whether our suspicion, that they were mostly used to refactor code, is true. We find that many of them were used to import the project or newer versions to git. The fact that such commits exist is unsurprising, considering that the investigated tools have been published before the release of git or before git was widespread. While their purpose might not be refactoring code, it does make sense to filter these commits, as they are not accurately depicting the development process of the

Table 5.3: Additional Information for Dataflow Analysis of Commits

Projects	Number of Active Commits	Probability for Commit to be part of dataflow CFI given that it is part of structural CFI
xz	1039	0.883
bzip2	37	0.889
gzip	194	1.000

project itself or of its features. Only two out of the eighteen commits we filtered in all projects, were actually used to refactor code, specifically moving contents between files or adapting the code to a newer version of C. For half of all exceptionally large commits we found no reason, neither refactoring code nor importing the project to git, to be left out of our analysis.

RQ2: Evaluation of Dataflow-based CFIs

Although dataflow-based CFIs only focus on dataflow in one direction, namely from commits to features, we examine properties for both commits and features and investigate them separately. Our results regarding the commits affecting features through dataflow are shown in figure 5.3 and in the tables 5.3, 5.4 and 5.5. Figure 5.4 displays our results for the features of a project and the commits that affect them through dataflow.

Proportion and Dependencies of Commits Affecting Features through Dataflow

At first, we examine the fraction of commits affecting features through dataflow among the active commits of a project. There, we also evaluate the effects and dependencies of commits part of structural CFIs on said fraction. Following this, we intend to determine the proportion of outside dataflow origins within the set of commits with dataflow interactions. This also allows us to quantify the commits whose interactions with features can only be discovered through our dataflow analysis. The number of active commits are shown in the first column of table 5.3 for each project. We determined xz to have the highest number of active commits at 1039, while bzip2 only has a tiny fraction of that at 37. In the first plot of figure 5.3, the bars colored in red show what percentage of commits interact with features through dataflow. We notice that the respective percentages are vastly different from project to project. The majority of gzip's active commits are part of dataflow-based CFIs at 62.9%. This percentage drops to about 27% for bzip2, with another large reduction to 11.3% for xz. This means that more than every second commit in gzip, roughly every fourth commit in bzip2 and every ninth commit in xz affects features through dataflow. The bars colored in grey display the fraction of active commits structurally interacting with features. In addition to the obvious fact of how often commits are used to implement features, this also bears important consequences for the fraction of commits with dataflow interactions. Comparing the two bar-types, we notice that the percentage of commits with structural interactions is lower than the percentage of commits with dataflow interactions for each project.

This phenomenon can be explained considering the values presented in the second column

Table 5.4: Commits exclusively/partially interacting with features only through dataflow

	Commit Count		Percentage of all active commits	
	Exclusive	Partial	Exclusive	Partial
xz	66	23	6.4	2.2
bzip2	2	6	5.4	16.4
gzip	31	54	16.0	27.9

of table 5.3. There, we show the probability for a commit to be part of any dataflow-based CFI, given that said commit is part of any structural CFI. We see that the probabilities are slightly below 90% for xz as well as bzip2 and even at 100% for gzip. This means that by only taking into account commits part of structural CFIs, we already encounter a lot of commits that affect features through dataflow. If we then consider the entire set of active commits, the number of commits with dataflow interactions will likely exceed the number of commits with structural interactions. This also implies that most, if not all, commits constituting the grey bar of a project in turn constitute parts of the red bar as well.

We explore the topic of different dataflow interaction types more thoroughly with our second plot, which focuses on the origin of dataflow. There, we have a closer look at the commits affecting features through dataflow and where their dataflow stems from. In the second plot of figure 5.3, we separate the commits with dataflow interactions into three categories based on dataflow origin. The first category, shown as the bar colored in blue, represents commits that only interact with features through outside dataflow. They make up the majority of commits for xz at 56.4%, but only 20% as well as 25.4% of commits for bzip2 and gzip respectively. The second category represents commits that interact with features through outside and inside dataflow. Logically, these commits affect at least two features through dataflow and only structurally interact with a subset of them. Surprisingly, they form the majority of commits for bzip2 and gzip at 60 and 44.3% respectively and around 20% of commits for xz. The proportion of commits that only interact with features through inside dataflow varies less between the projects. gzip has the highest percentage at 31.3%, bzip2 the lowest at 20% with xz inbetween the two at 23.9%.

Combining the data presented in both plots, we calculated the number and percentage of active commits in a project whose interactions with features were exclusively discovered by our dataflow analysis. The respective values are shown in table 5.4 on the left-hand side of each column. With a count of 66, most commits were revealed for xz, although its percentage remains comparatively small due to its high number of active commits. gzip encompasses by far the highest percentage at 16% with a commit count of 31. On the right-hand side of the columns in table 5.4 we display the number and percentage of commits where only a subset of their interactions with features was discovered by our dataflow analysis. These commits are most common in gzip with a count of 54 making up more than every fourth active commit at 27.9%. Similarly to the overall percentage of commits with dataflow interactions, they are less common in bzip2 at 16.4% and least common in xz at 2.2%.

We notice that commits interacting with features through inside dataflow, have a high chance to also interact with other features through outside dataflow. By definition, all commits

Table 5.5: Relating Inside Dataflow to Outside Dataflow

Projects	Probability for Commits Part of Structural CFIs to Affect other Features Through Outside DF	Same Probability Given Any Active Commit
xz	0.417	0.086
bzip2	0.667	0.216
gzip	0.596	0.381

interacting with features through inside dataflow also structurally interact with them. From previous explanations, we also know these commits are almost identical to the entire set of commits with structural interactions. It follows that commits part of structural CFIs must also have a high, albeit slightly lower, chance to affect other features through outside dataflow. We compare this probability to the probability of any active commit to interact with features through outside dataflow in table 5.5. The strongest contrast occurs for xz, where 41.7% of commits structurally interacting with features, interact with other features through outside dataflow, making them 4.8 times more likely to do so compared to any active commit of xz at 8.6%. The contrast is slightly weaker for bzip2 with a 3.1 times higher likelihood and according percentages of 66.7 and 21.6% respectively. gzip has the weakest contrast at 1.5 and respective probabilities of 59.6 and 38.1%.

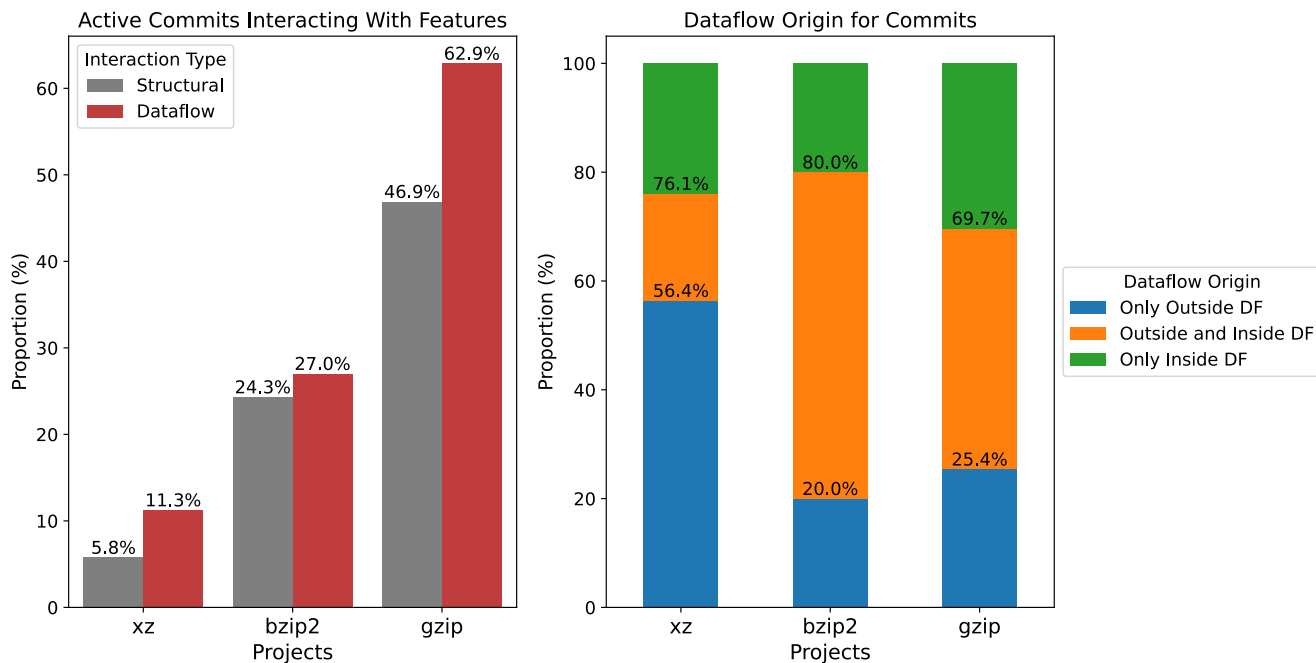


Figure 5.3: Proportional Dataflow-Plot for Commits

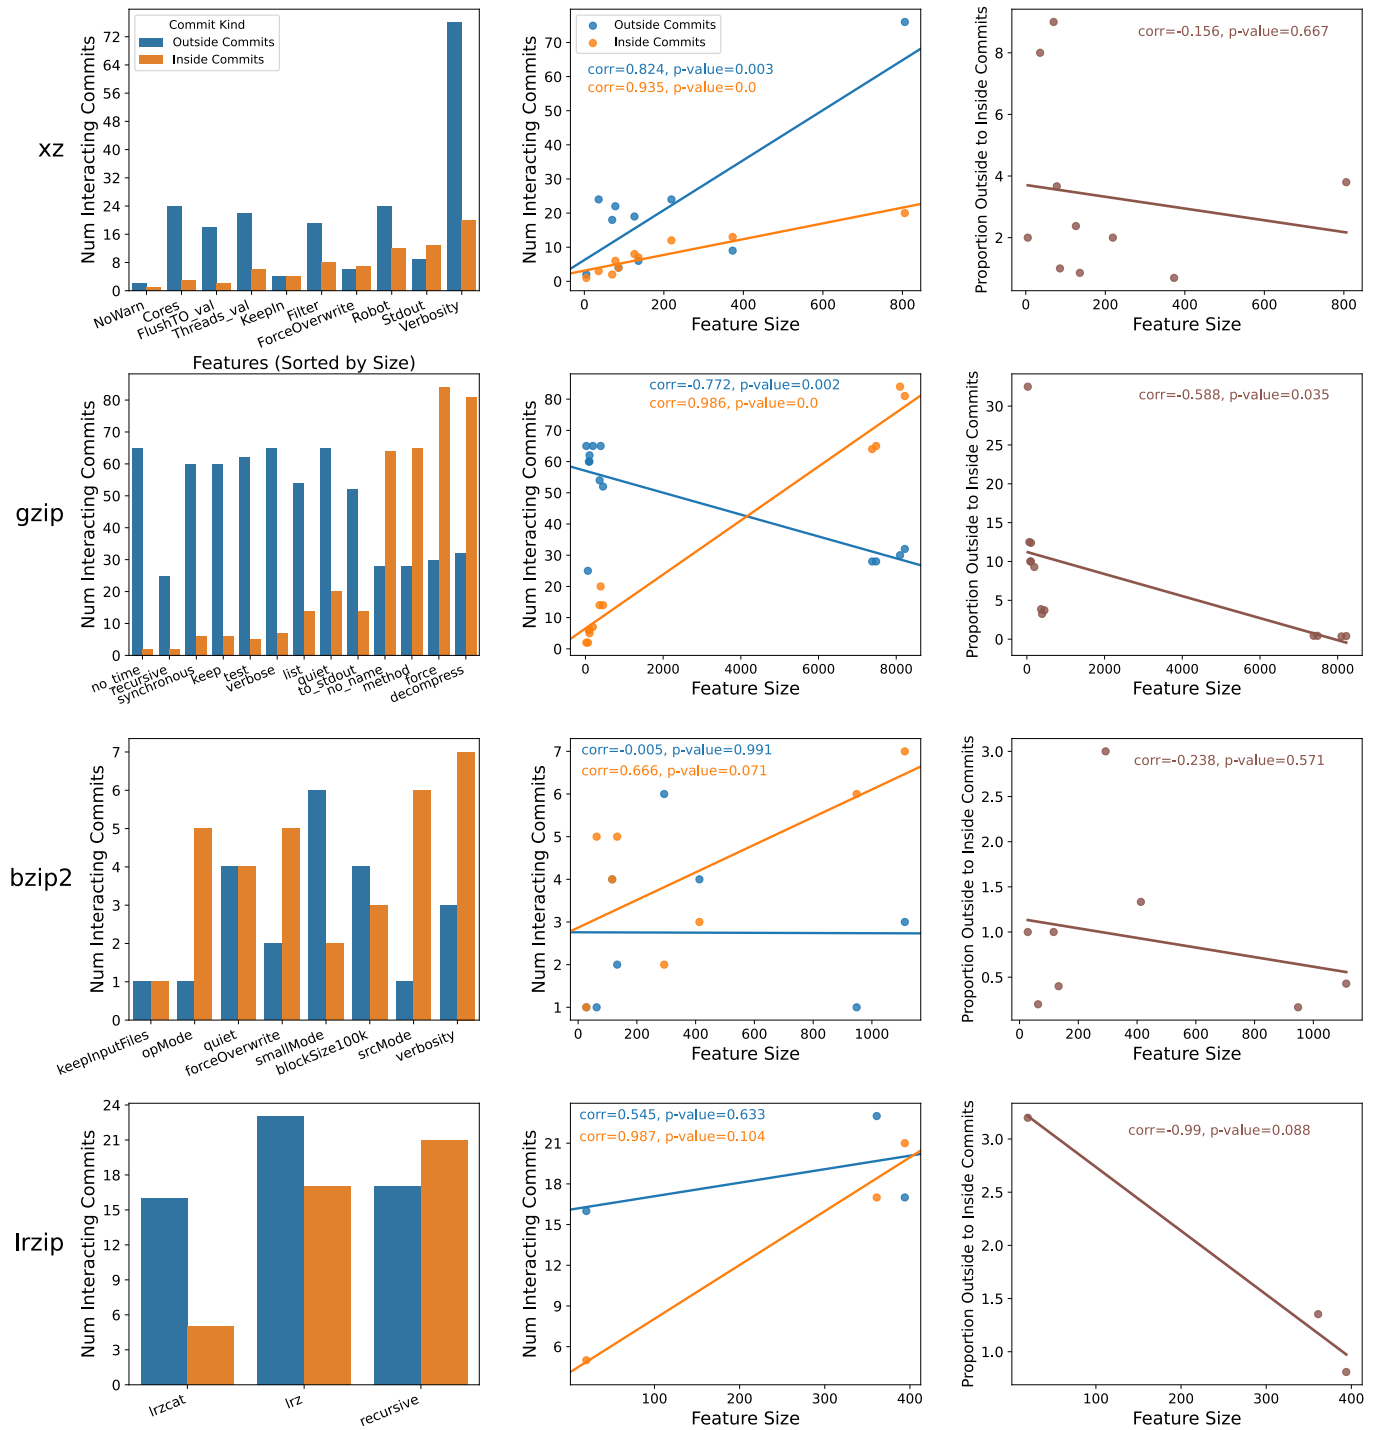


Figure 5.4: Proportional Dataflow-Plot for Features

Understanding Features and the Commits Affecting them Through Dataflow

Besides examining the ratio of outside and inside commits affecting a feature through dataflow, we also relate the two metrics to the size of a feature. Furthermore, we want to put our thesis, that the proportion of outside to inside commits of a feature decreases as its size increases, to the test. Our results are shown in figure 5.4 in three plots for each of our investigated projects. In column one of the respective figure, we display the number of outside and inside commits for each feature in a bar plot. We note that the majority of features in xz, GZIP and LRZIP are affected by more outside than inside commits through dataflow. For xz and GZIP this is underlined in median-values of 18.5 and 54 for outside commits and median-values of 6.5 and 14 for inside commits among their features. For features of xz and GZIP with more outside than inside commits, we also notice very high ratios of outside to inside commits. They range from 2 to 9 in xz with a mean of 4.4 and from 3 to 32 in GZIP with a mean of 10.8. For features with more inside than outside commits affecting them through dataflow, we determined much lower ratios and according ranges of inside to outside commits. They range from 1.17 to 1.44 in xz with a mean of 1.3 and from 2.3 to 2.8 in GZIP with a mean of 2.5. At least for these two projects, the number of outside commits is often not only slightly higher than the number of inside commits, but instead exceeds it by many magnitudes. The features of BZIP2 exhibit a different general behavior than the features of the projects discussed above. Specifically, 6 out of 8 features encompass more or at least the same number of inside than outside affecting them through dataflow. The range and maximum number of interacting commits is also much lower compared to the other examined projects.

In the second column of figure 5.4, we compare the number of inside and the number of outside commits of a feature with its size in a regression plot. As suspected during the formulation of RQ2, we find a strong positive correlation regarding the inside commits of features across all projects. Especially, the correlation coefficients for xz, GZIP and LRZIP range close to 1 with a minimum value of 0.935. xz and GZIP encompass p-values of less than 10^{-4} supporting a statistically significant correlation. The correlations regarding the outside commits affecting features through dataflow are vastly different from project to project. xz is the only project with a strong positive correlation with a coefficient of 0.824. A low p-value of 0.003 outside of our rejection interval makes the data seem statistically significant at first. However, the dots of its regression plot are scattered rather randomly and the number of interacting commits generally does not increase with rising feature-size. We determined the Verbosity-feature to be the main reason for our calculated strong positive correlation. This becomes clear when we disregard Verbosity in our analysis, which results in the coefficient dropping close to 0 and the p-value rising to 1. GZIP encompasses a negative correlation coefficient of -0.772 and a p-value of 0.002 justifying the existence of significant statistical evidence. The falling linear regression line is mainly formed by the two existing clusters, where one consists of four very large and the other cluster of nine relatively small features. The former cluster consists of features that each interact with around 30 outside commits through dataflow. The cluster consisting of features with small sizes encompasses higher values between 50 and 65. The only feature of GZIP that cannot be assigned to a cluster is the recursive-feature, which rather fits into the first cluster according to its number of outside commits and into the second cluster according to its size. In BZIP2, the number of outside commits affecting a feature through dataflow and its size are un-correlated. A correlation

coefficient of 0.545 and a high p-value of 0.633 do not provide significant statistical evidence for the existence of a correlation in LRZIP.

In the third column of figure 5.4, we finally relate the size of a feature to its proportion of outside to inside commits. At first, we notice that the calculated correlation coefficients are negative across all projects. The fact that the coefficients of xz and BZIP2 are close to 0 and the p-values correspondingly high is illustrated by the rather random distribution of dots in their respective regression plots. Their data should therefore not be regarded as statistically significant. Compared to the aforementioned projects, GZIP's p-value of 0.035 is very low, while it also has a lower correlation coefficient of -0.588. Similarly to GZIP's previous plots, we can identify two clusters. A cluster of dots is again formed by the four largest features, encompassing proportions of outside to inside commits slightly below 0.5. The other cluster is located at the beginning of the x-axis with a wide spread across the y-axis, indicating its low feature-size and a wide range of proportions from 3 to 32. The coefficient of LRZIP is by far the smallest among all projects at -0.99, but the small number of features results in a p-value of 0.088. It follows that none of the determined negative correlations are statistically significant, as their p-values fall into our rejection interval of 97.5.

RQ3: Evaluation of Author Interactions

- the project xz has mostly been developed by a single author, which is reflected in our results
- as can be seen in figure 5.5, all except one feature have exclusively been implemented by the same author
- we can conclude that it also is the developer affecting features through outside dataflow, since no feature has unique dataflow authors
- it follows that the linear regression relating the size of a feature to the number of its interacting authors shows no significant correlation
- BZIP2 inconclusive data, no correlation to be found
- LRZIP more authors of features, but no general trend to be found due to lack of features
- GZIP, on the other hand, provides us with much more comprehensive results
- it can immediately be seen that every feature has a large number of interacting authors
- the overall number of authors that interact with a feature in any way can be computed by the sum of its structural and its unique dataflow authors
- this means that the lowest number of interacting authors is 7 for the recursive-feature and the highest number 14 for the force-feature
- on average, a feature has more dataflow than structural authors with respective averages of 9.1 and 5.7
- the average number of unique dataflow authors is 5.9, which means that our dataflow analysis discovers many additional interacting authors in GZIP
- the number of structurally interacting authors, i.e. authors implementing a feature, range from 2 to 13 between the features of GZIP
- there is less variance in the number of a feature's outside dataflow authors ranging from 7 to 12
- features with relatively few structural authors, have a comparatively high number of dataflow

authors

- features with many structural authors, starting from the NO_NAME-feature, have less outside dataflow authors
- at the same time, the number of their unique dataflow authors is rather low, implying that most of them also structurally interact with the according features
- the fact that the group of features with few structural authors has much more dataflow authors conditions them to have many unique dataflow authors
- still, more than 80% of their structural authors also interact with them through outside dataflow
- the same cannot be said for the group of features with many structural authors, where this is the case for less than 45% of them
- overall, there is a 57% likelihood for the structural author of any feature to also interact with the feature through outside dataflow
- similarly to the bar-plot of GZIP, we notice two clusters when examining its regression plots in figure 5.5
- the first cluster is made up of features with few structural authors, whose respective sizes range between 50 and 500
- the second cluster consists of the four features who all have over 10 structural authors and feature-sizes ranging from 5500 to 6500
- it is unsurprising that we compute a strong positive correlation between the size of a feature and its structurally interacting developers
- the respective correlation coefficient is 0.976, while the p-value is smaller than 10^{-4}
- regarding the dataflow authors of features, we note two similar clusters that consist of the same features of smaller and large sizes respectively
- now, the smaller features have slightly more dataflow authors centering around 10 than the much larger features with an average of 8
- this shows itself in a negative correlation coefficient of -0.746 and a p-value of 0.003 falling out of our rejection range
-




Figure 5.5: Author CFI Plot

5.2 DISCUSSION

RQ1: Structural CFIs

Generally, we have seen our expectations for possible results of RQ1 mostly confirmed. Here, we briefly formulate our most important results, which we discuss in more detail in the following sub-sections. Firstly, we found sufficient evidence to support the notion that the number of commits used to develop a feature is strongly positively correlated with the size of the code implementing its functionality. Secondly, we have high confidence that most commits, within our investigated projects, only changed or implemented functionality of one feature.

Patterns around Feature Development

The features within the investigated projects exhibit a wide distribution of structurally interacting commits and sizes. We consider this as a clear indication that the number of commits as well as the extent of source-code used to implement a feature varies a lot between them. The range of least to most interacting commits and lowest to highest feature-size is hereby dependent on the specific project.

Determining the average nesting degree of features within a project showed us that, especially in the case of GZIP, feature-nesting is more common than initially expected. Most of the computed sizes of GZIP's respective features are made up of instructions that are part of multiple feature-regions besides its own. This makes it difficult to clearly interpret the data, as the number of instructions that actually implement a feature can differ greatly from our observed size. This phenomenon also bears consequences on the way commits are used in feature development and the CFIs resulting from this. If a developer wants to specifically change a feature nested inside other features, the according commit will necessarily produce structural CFIs with several features. This also casts some doubt how well the structural interactions of a feature with a high nesting degree accurately represent its development. A concrete example of this are the four, by far largest, features of GZIP. In our evaluation of RQ1, we have already explained that they share the vast majority of instructions constituting their size, as well as the commits that structurally interact with them. In these cases, it would be extremely helpful to be able to decide whether the respective instructions and commits are actually implementing commits and instructions of a particular feature. While solving this problem goes beyond the scope of this work, we will address the subject and possible solutions again in the [Future Work](#) section of this thesis.

We employed the pearson correlation coefficient to investigate the linear relationship between the size of a feature and the number of commits structurally interacting with it. While we found high correlation coefficients of 0.9 and above, the p-values of LRZIP and BZIP2 were not low enough to justify the data to be statistically significant. We know that p-values are highly dependent on the number of datapoints, that is, the number of features of a project. Logically, projects with fewer features have less datapoints resulting in their p-values being quite high and falling into our rejection interval. With 10 and 14 features respectively, xz and GZIP have by far the most datapoints and according p-values of less than 10^{-4} giving us most significant evidence for a strong positive correlation. Due to the above-mentioned problems caused by

GZIP's high overall nesting degree, the significance of the found strong positive correlation is somewhat weakened. On the other hand xz has a relatively low overall nesting degree, but still shows the same strong positive correlation. Therefore, we are still able to confirm our original thesis of a strong dependency between the size of the implementing code and the number of implementing commits of a feature.

A possibility to circumvent the problem of too few datapoints could be to combine features of all projects into one big dataset. Here, we are faced with the issue that the way commits are used among projects might differ drastically. In our case, the number of instructions per structurally interacting commit of a feature varies greatly from project to project. The respective values are 17 for xz, 48 for GZIP, 87 for BZIP2 and 13 for LRZIP. This means that features of a similar size have, on average, vastly different numbers of structurally interacting commits across different projects. Upon these considerations, we are of the opinion that it makes more sense to test correlations within projects as we have done in this work.

Usage of Commits in Feature Development

In section 3.5, we explained that, due to feature-nesting, the number of features a commit structurally interacts with can be a faulty estimate of how many features a commit usually changes. To achieve a more accurate measure, we introduced the concept of feature-related concerns of a commit (8). There, we do not look at the structural CFIs of a commit separately, but rather focus on the entirety of them and the instructions they stem from. As long as the CFIs occur in the same feature-related location, i.e. in instructions belonging to regions of the same set of features, we assume that only one specific feature was implemented by the commit. A concrete example of why the feature-related concerns can be a more accurate estimate for the number of changed features is shown in listing 3.5. Excluding BZIP2, roughly 70% of commits only have a single feature-related concern in the remaining projects. Despite this, GZIP's average is quite high at 2.4, in comparison to xz and LRZIP with 1.48 and 1.3 respectively. It decreases to 1.55 after filtering large commits, such as commit 33ae4134cc with an exceptionally high number of 47 feature-related concerns. Its purpose was to import the original gzip.c-file into the repository, which justifies the exclusion of the commit. Besides that, GZIP only encompasses 14 features making it impossible for a commit to change 15 or more features. The number of feature-related concerns being higher than the overall number of features also occurs for commits of BZIP2 and LRZIP. It follows that the concerns of commits can sometimes overestimate the number of changed features. Naturally, this is most likely for commits structurally interacting with many features that also have high nesting degrees. The code affected by these commits can include more combinations of different sets of feature regions than features present in the project. For the respective commits, the number of feature-related concerns should rather be reduced to the total number of features they interact with. To summarize our results, most commits of xz, GZIP and LRZIP have only one feature-related concern with an average of less than 1.6 after filtering large commits. Given that the feature-related concerns accurately reflect the number of changed features for most commits, we can confirm our initial thesis. Thus, the majority of commits that structurally interact with features are only used to change functionality of one feature.

In contrast to three projects discussed above, BZIP2 has shown to be a project whose results need to be treated more carefully. Only 9 commits structurally interact with its features,

encompassing much less datapoints compared to the other investigated projects. 6 out of 9 commits were used to import the initial project to git and iteratively update the code to newer versions. These commits are not exemplary of the common development process of a project and its features in a git repository. The actual development of the tool in the git repository began after `BZIP2` was finally updated to version 1.0.6. The remaining 3 commits introduce minor changes and small bug fixes to the respective features they interact with. This means that `BZIP2`'s features were largely already implemented when the actual development within the repository started. Thus, the average number of concerns of a commit as well as their overall distribution are not good representatives of other projects developed in git.

RQ2: Dataflow-based CFIs

Proportion and Dependencies of Commits Affecting Features through Dataflow

We have seen that the proportion of commits part of dataflow-based CFIs varies a lot between our investigated projects. While only 11% of commits in `xz` affect features through dataflow, the majority of commits do so in `GZIP`. One major reason for the discussed differences is the fact that the proportion of structural commits acts as the lowest bound for the proportion of commits with dataflow interactions. As 90-100% of commits part of structural CFIs are also part of dataflow-based CFIs, the percentage of commits part of dataflow-based CFIs never falls below the percentage of commits part of structural interactions. The effects of said lowest bound can be recognized by the fact that a project's rank among the dataflow proportions of commits reflects its rank among the structural proportions. That is also because the percentages of commits, whose interactions with features were only discovered by our dataflow analysis, are the same or only a fraction of the percentages of those with structural interactions. Similar strong differences in the structural proportions of commits can thus also be found in the dataflow proportions. Including structural CFIs in the evaluation of our dataflow-based CFIs was definitely necessary to properly categorize their proportion and the new information they could add to the analysis of a project. A substantial share of the examined dataflow interactions stem from our previously investigated set of structural interactions. As we have determined structural CFIs to heavily coincide with dataflow-based CFIs, we already know of this set of dataflow interactions prior to our dataflow analysis. Nevertheless, we fully believe that our dataflow analysis is an important part of understanding commit-feature interactions within a project. For `xz`, the majority of commits interacting with features were exclusively revealed by our dataflow analysis. For 8, 21 and 44% of active commits in `xz`, `BZIP2` and `GZIP`, it revealed additional interactions with features that we could not discover with our structural analysis. This means that the data of, on first sight unrelated, commits and features can be highly connected. Besides that, we found that many commits interacting with features both structurally and through dataflow, also interact with other features only through dataflow. This lead us to examine whether commits structurally interacting with features are more likely to affect features through outside dataflow in comparison to any active commit. We found sufficient evidence to support the notion that this is indeed the case, including an almost 5-times higher probability in `xz`. This suggests that data, changed or allocated by a commit in one feature, is often used within a different feature again. Therefore, a developer changing code of a feature should be aware that data of said code has an increased

probability to affect the functionality of other features. In the following section, we will also seek to underline this notion by determining how often commits affecting a feature through outside dataflow, structurally interact with other features.

Lastly, we discuss the extent of feature-code in a project and its consequences on the proportion of commits with structural as well as dataflow interactions. For xz and GZIP, we have confirmed that a feature's number of structurally interacting commits is strongly positively correlated with its size. It follows that a higher extent of feature-code implies a higher proportion of commits with structural interactions. As the latter proportion acts as a lower bound for the proportion of commits with dataflow interactions, we assume the extent of feature-code to also have a strong influence on said proportion. At the same time less feature-code and a higher number of active could elavate the share of commits affecting features through outside dataflow. This is due to a lower share of structurally interacting commits and therefore less inside dataflow as well as a higher number of active commits located outside of feature-regions leading to more possible outside dataflow. We decided to compute the extent of feature-code in a project as the fraction of all instructions that are part of feature regions. Thus, we determined the extent of feature-code to be 17.2% in xz and 57.6% in GZIP. In our evaluation we have also seen that xz encompasses five times as many active commits than GZIP. We find that our observations in RQ2 roughly match these values considering our previous explanations. Accordingly, xz has a higher share of commits affecting features through outside dataflow but a substantially lower share of active commits part of dataflow interactions. It should be noted however that the extent of feature-code does not have an exaction relation of the proportion of commits with structural interactions. From xz's 17% extent of feature-code, we would expect a significantly higher number than the 5.8% of commits part of structural CFIs for example.

Understanding Features and the Commits Affecting them Through Dataflow

- especially in the case of xz and GZIP, we have found significant evidence to support the notion that features are generally affected by more outside than inside commits through dataflow
 - this means that most detected dataflow interactions of features originate outside of its regions and subsequently flows into them
 - the ratio of outside to inside commits is often extremely skewed towards outside commits, while the opposite is rarely the case and if so, to a much smaller degree
 - one reason for this is that even features with little size are affected by many outside commits and only few inside commits
 - that is because the number of a feature's inside commits is strongly positively correlated with its size, while the same is not the case for a feature's outside commits
 - specifically, we did not find evidence to suggest that the size of a feature and the number of outside commits affecting it through dataflow are correlated in some way or another
 - the positive correlation found for xz is mainly driven by its largest feature, where the correlation coefficient subsides to 0 without it
 - GZIP shows a statistically significant negative correlation, BZIP2 no correlation and LRZIP a statisitcally insignificant positive correlation
-
- proportion outside to inside commits corr with size

- outside commits also structurally interact with other features => suggest dataflow occurs within features
- correlation of gzip mainly driven by two clusters discussed in our evaluation => in eval of RQ1 we also found out that the commits the features of the cluster with large size interact with are largely identical => we found the same to be true for the outside commits of cluster with small sizes and for the outside and inside commits of the other cluster => due to feature-nesting => makes the data perhaps less significant, as the data of many commits could flow into instructions that do not implement functionality of the features the commits interact with

5.3 THREATS TO VALIDITY

In this section, discuss the threats to internal and external validity.

RELATED WORK

Interactions between features [2, 3] and interactions between commits [6] have already been used to answer many research questions surrounding software projects. However, investigating feature interactions has been around for a long time, while examining commit interactions is a more recent phenomenon.

In an article published in 2023, Sattler et al. [6] analysed several open-source projects with their novel approach, SEAL. SEAL merges low-level data-flow with high-level repository information in the form of commit interactions. The paper shows the importance of a combination of low-level program analysis and high-level repository mining techniques by discussing research problems that neither analysis can answer on its own. For example SEAL is able to detect commits that are central in the dependency structure of a program[6]. This was used to identify small commits affecting central code that would normally not be considered impactful to a program. Furthermore, they investigated author interactions at a dataflow level with the help of commit interactions. Thus, they can identify interactions between developers that cannot be detected by a purely syntactical approach. They found that, especially in smaller projects, there often exists one main developer authoring the majority of commits[6] and thus accounting for most author interactions logically. It was also explained how SEAL makes it possible to relate occurrences of bad programming practices to developers. This is accomplished by SEAL enriching program analyses with computed repository information. Lillack et al. [3] first implemented the automatic detection of features inside programs by tracking their load-time configuration options along program flow. In our research, we also focus on features who are configured via configuration options, which we call configuration variables. Their analysis tool Lotrack can detect which configuration options must be activated in order for certain code segments, implementing a feature’s functionality, to be executed. They evaluated Lotrack on numerous real-world Android and Java applications and observed a high accuracy for the predicted code execution constraints[3].

Referencing this paper Kolesnikov et al. [2] published a case study on the relation of external and internal feature interactions. Internal feature interactions are control-flow feature interactions that can be detected through static program analysis as mentioned above. They concluded that considering internal feature interactions could potentially help predict external, performance feature interactions[2].

CONCLUDING REMARKS

7.1 CONCLUSION

7.2 FUTURE WORK

APPENDIX

This is the Appendix. Add further sections for your appendices here.

BIBLIOGRAPHY

- [1] Christopher Hundhausen, Adam Carter, Phillip Conrad, Ahsun Tariq, and Olusola Adesope. "Evaluating Commit, Issue and Product Quality in Team Software Development Projects." In: SIGCSE '21. Virtual Event, USA: Association for Computing Machinery, 2021, pp. 108–114. ISBN: 9781450380621. DOI: [10.1145/3408877.3432362](https://doi.org/10.1145/3408877.3432362). URL: <https://doi.org/10.1145/3408877.3432362>.
- [2] Sergiy Kolesnikov, Norbert Siegmund, Christian Kästner, and Sven Apel. "On the relation of external and internal feature interactions: A case study." In: *arXiv preprint arXiv:1712.07440* (2017).
- [3] Max Lillack, Christian Kästner, and Eric Bodden. "Tracking load-time configuration options." In: *Proceedings of the 29th ACM/IEEE international conference on Automated software engineering*. 2014, pp. 445–456.
- [4] Florian Sattler. "Understanding Variability in Space and Time." To appear. dissertation. Saarland University, 2023.
- [5] Florian Sattler. *VaRA is an analysis framework that enables users to build static and dynamic analyses for analyzing high-level concepts using advanced compiler and analysis technology in the background*. <https://vara.readthedocs.io/en/vara-dev/> [Accessed: (24.05.2023)]. 2023.
- [6] Florian Sattler, Sebastian Böhm, Philipp Dominik Schubert, Norbert Siegmund, and Sven Apel. *SEAL: Integrating Program Analysis and Repository Mining*. 2023.